

Project: Building Ingest-Transform-Load Streaming Pipeline

link for short youtube video:

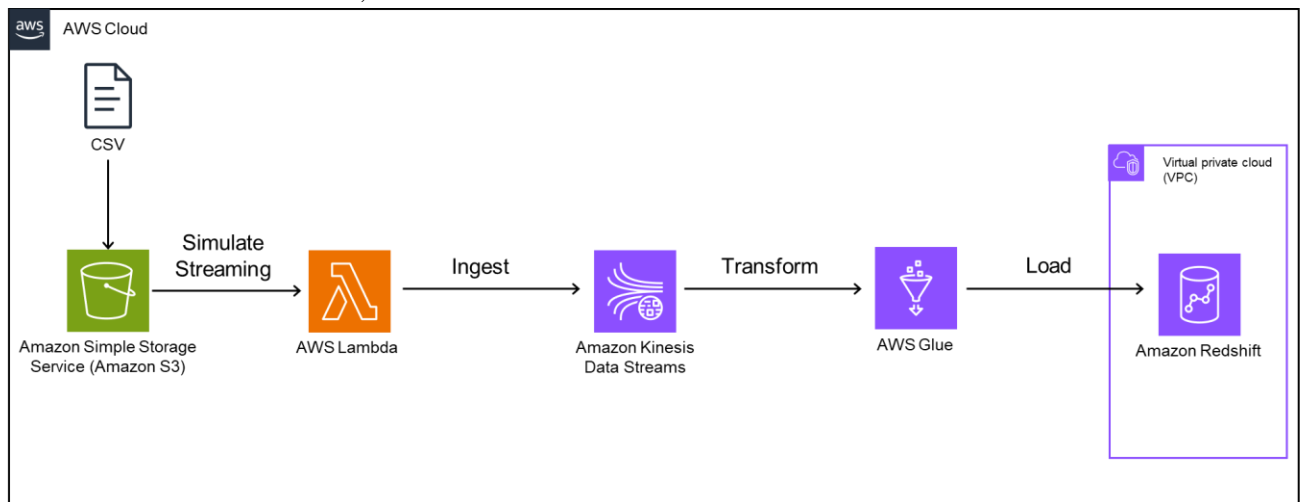
https://youtu.be/B6mv_ffa0Pc

link for long youtube videos:

<https://youtu.be/p9z2FUVVbIU>

Project Overview

Kinesis Data Stream is a service that offers to process data ingestion that is high in volume and velocity from various data sources and deliver to the desired end point. Kinesis Data Stream also offers fault-tolerant and data distribution across shards based on a partition key, which makes it suitable to be used in ETL process for streaming data. In this demo, we are going to demonstrate the use of Kinesis Data Stream in the ingest-transform-load streaming pipeline of three baseball players three months game statistics, including Shohei Ohtani, Betts Mookie, and Freddie Freeman. The statistics CSV file will be saved in S3 bucket. Lambda is used to stream each record to the Kinesis Data Streaming. After ingestion, the data will be transformed in Glue by selecting only the important statistics. After that, data will be loaded into a warehouse, Amazon redshift.



Create a Virtual Private Cloud and Private Subnets

The VPC and subnets are created to only host Amazon Redshift

- Go to VPC service in AWS
- Click on Create VPC
- Name it ITL-vpc
- Choose IPv4 CIDR manual input
- For IPv4 CIDR, use 10.0.0.0/24 for amount of 4,096 IPs
- Select No IPv6 CIDR block
- Click create VPC

Your VPC list should look as the following:

[VPC](#) > [Your VPCs](#) > vpc-070fdfc40b96963e9

vpc-070fdfc40b96963e9 / ITL-vpc Actions ▼

Details [Info](#)

VPC ID vpc-070fdfc40b96963e9	State Available	Block Public Access Off	DNS hostnames Disabled
DNS resolution Enabled	Tenancy Default	DHCP option set dopt-0542a3f2947465eb3	Main route table rtb-05c76cab9fa74ae39
Main network ACL acl-066ae32f78e38b5db	Default VPC No	IPv4 CIDR 10.0.0.0/24	IPv6 pool –
IPv6 CIDR –	Network Address Usage metrics Disabled	Route 53 Resolver DNS Firewall rule groups –	Owner ID 124355677069

[Resource map](#) | [CIDRs](#) | [Flow logs](#) | [Tags](#) | [Integrations](#)

Resource map [Info](#)

VPC [Show details](#)
Your AWS virtual network

ITL-vpc

Subnets (0)
Subnets within this VPC

Route tables (1)
Route network traffic to resou

rtb-05c76cab9fa74ae39

- Go to subnets
- Click create subnet
- Name subnet us-west-1a-subnet
- Choose the Availability Zone in your area. I choose us-west-1a
- Choose 10.0.0.0/24 for IPv4 VPC CIDR block
- Choose 10.0.0.0/25 for IPv4 subnet CIDR block with 128 IP addresses
- Create subnet

Your subnet should look as the following:

subnet-0bcc4f75df83d370a / us-west-1a-subnet

Actions ▼

Details

Subnet ID ☐ subnet-0bcc4f75df83d370a	Subnet ARN ☐ arn:aws:ec2:us-west-1:124355677069:subnet/subnet-0bcc4f75df83d370a	State ✔ Available	Block Public Access ⊖ Off
IPv4 CIDR ☐ 10.0.0.0/25	Available IPv4 addresses ☐ 123	IPv6 CIDR -	IPv6 CIDR association ID -
Availability Zone ☐ us-west-1a	Availability Zone ID ☐ usw1-az1	VPC vpc-070fd40b96963e9 ITL-vpc	Route table rtb-05c76cab9fa74ae39
Network ACL -	Default subnet No	Auto-assign public IPv4 address No	Auto-assign IPv6 address No
Auto-assign customer-owned IPv4 address No	Customer-owned IPv4 pool -	Outpost ID -	IPv4 CIDR reservations -
IPv6 CIDR reservations -	IPv6-only No	Hostname type IP name	Resource name DNS A record Disabled
Resource name DNS AAAA record Disabled	DNS64 Disabled	Owner ☐ 124355677069	

- Create another subnet name us-west-1c-subnet
- Choose the Availability Zone in your area. I choose us-west-1c
- Choose 10.0.0.0/24 for IPv4 VPC CIDR block
- Choose 10.0.0.128/25 for IPv4 subnet CIDR block with 128 IP address

Your subnet should look as the following:

subnet-0ce82e6faa060ecbd / us-west-1c-subnet

Actions ▼

Details

Subnet ID ☐ subnet-0ce82e6faa060ecbd	Subnet ARN ☐ arn:aws:ec2:us-west-1:124355677069:subnet/subnet-0ce82e6faa060ecbd	State ✔ Available	Block Public Access ⊖ Off
IPv4 CIDR ☐ 10.0.0.128/25	Available IPv4 addresses ☐ 123	IPv6 CIDR -	IPv6 CIDR association ID -
Availability Zone ☐ us-west-1c	Availability Zone ID ☐ usw1-az3	VPC vpc-070fd40b96963e9 ITL-vpc	Route table rtb-05c76cab9fa74ae39
Network ACL -	Default subnet No	Auto-assign public IPv4 address No	Auto-assign IPv6 address No
Auto-assign customer-owned IPv4 address No	Customer-owned IPv4 pool -	Outpost ID -	IPv4 CIDR reservations -
IPv6 CIDR reservations -	IPv6-only No	Hostname type IP name	Resource name DNS A record Disabled
Resource name DNS AAAA record Disabled	DNS64 Disabled	Owner ☐ 124355677069	

- Check in route tables, there should be a default route tables created

- Name it private-subnets-route
- 10.0.0.0/24 allows routing inside the VPC

VPC > Route tables > rtb-05c76cab9fa74ae39

rtb-05c76cab9fa74ae39 / private-subnets-route Actions ▼

Details Info

Route table ID rtb-05c76cab9fa74ae39	Main Yes	Explicit subnet associations -	Edge associations -
VPC vpc-070fd40b96963e9 ITL-vpc	Owner ID 124355677069		

Routes Subnet associations Edge associations Route propagation Tags

Routes (1) Both ▼ Edit routes

< 1 > ⚙

Destination ▼	Target ▼	Status ▼	Propagated ▼
10.0.0.0/24	local	Active	No

Create IAM role for the demo

- Go to IAM and choose roles
- Click create role
- Choose AWS account for select trusted entity and choose this account
- Attach the following policies:
 - AmazonS3ObjectLambdaExecutionRolePolicy for S3 to interact with Lambda
 - AWS Lambda_fullAccess
 - AWSS3FullAccess
 - AWSLambdaKinesisExecutionRole
 - AmazonKinesisFullAccess
 - AmazonRedshiftFullAccess
 - AWSGlueConsoleFullAccess
- Name Role Name data-engineer
- Create role

You should be able to see your role as this:

data-engineer [Info](#)

[Delete](#)

Summary

[Edit](#)

Creation date

December 03, 2024, 21:29 (UTC-08:00)

ARN

[arn:aws:iam::124355677069:role/data-engineer](#)

Link to switch roles in console

<https://signin.aws.amazon.com/switchrole?roleName=data-engineer&account=protprommart>

Last activity

-

Maximum session duration

8 hours

Permissions

Trust relationships

Tags

Last Accessed

Revoke sessions

Permissions policies (6) [Info](#)

[Simulate](#)[Remove](#)[Add permissions](#)

You can attach up to 10 managed policies.

Filter by Type			
<input type="text" value="Search"/>			
All types			
< 1 > ⚙			
<input type="checkbox"/>	Policy name	Type	Attached entities
<input type="checkbox"/>	AmazonRedshiftFullAccess	AWS managed	1
<input type="checkbox"/>	AmazonS3FullAccess	AWS managed	1
<input type="checkbox"/>	AmazonS3ObjectLambdaEx...	AWS managed	1
<input type="checkbox"/>	AWSGlueConsoleFullAccess	AWS managed	1
<input type="checkbox"/>	AWSLambda_FullAccess	AWS managed	1
<input type="checkbox"/>	AWSLambdaKinesisExecuti...	AWS managed	1

- Once done, go to the trust relationship tab of the data engineer role we created.
- We are going to add services that we want to use this role with including redshift, lambda, glue, kinesis.
- In the json documents, use the following json to establish service trust

Trusted entities

[Edit trust policy](#)

Entities that can assume this role under specified conditions.

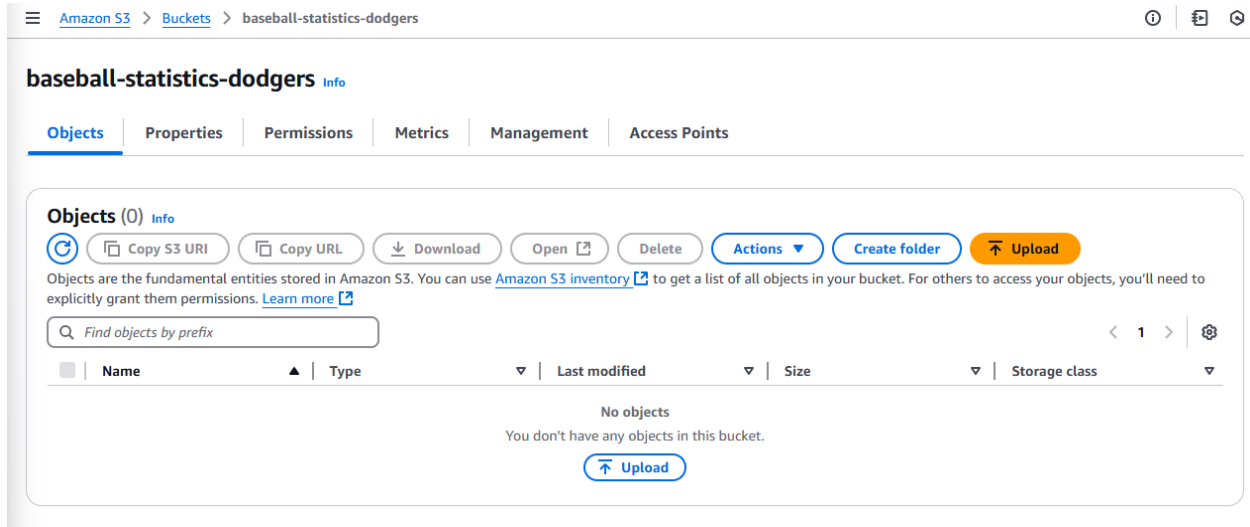
```
1 {
2   "Version": "2012-10-17",
3   "Statement": [
4     {
5       "Effect": "Allow",
6       "Principal": {
7         "Service": [
8           "redshift.amazonaws.com",
9           "glue.amazonaws.com",
10          "lambda.amazonaws.com",
11          "kinesis.amazonaws.com"
12        ]
13      },
14      "Action": "sts:AssumeRole"
15    }
16  ]
17 }
```

Create a S3 bucket

- Go to Amazon S3 bucket in AWS
- Click on Create bucket
- Name the bucket baseball-statistics-dodgers

- Choose ACLs disabled for object ownership
- Select block all public access
- Choose disable bucket Versioning
- Create bucket

You should have a blank bucket like this:



Input Baseball Statistics into S3 bucket

I extracted game log statistics from March to May 2024 of LA Dodgers players Shohei Ohtani, Betts Mookie, and Freddie Freeman from LA Dodgers websites:

<https://www.mlb.com/player/freddie-freeman-518692?season=2024&team=119&stats=gamelogs-r-hitting-mlb&year=2024>

<https://www.mlb.com/player/shohei-ohtani-660271?stats=gamelogs-r-hitting-mlb&year=2024>

<https://www.mlb.com/player/mookie-betts-605141?season=2024&team=119&stats=gamelogs-r-hitting-mlb&year=2024>

The csv file of the game log statistics from March to May 2024 can be found in my github repository link:

https://github.com/Protprommart/LA-Dodgers-Cloud-ETL/blob/main/baseball_players_march_to_may.csv










- Download csv file from the github link
- Upload this file into the baseball-statistics-dodgers S3 bucket

You should be able to see the file in the bucket like this:

baseball-statistics-dodgers info


Objects | Properties | Permissions | Metrics | Management | Access Points

Objects (1) Info





  Copy S3 URI  Copy URL  Download  Open  Delete  Actions  Create folder  Upload

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

< 1 > 

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	 baseball_players_march_to_May.csv	csv	December 3, 2024, 21:07:38 (UTC-08:00)	13.5 KB	Standard










- Create another S3 bucket called clean-baseball-statistics with the same instruction from before. This is going to be a bucket where we put in our transformed dataset after running the data with visual ETL. You should have the bucket as follow:

 [Amazon S3](#) > [Buckets](#) > clean-baseball-statistics   


clean-baseball-statistics Info

Objects | Properties | Permissions | Metrics | Management | Access Points

Objects (0) Info


  Copy S3 URI  Copy URL  Download  Open  Delete  Actions  Create folder  Upload

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

< 1 > 

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
--------------------------	------	------	---------------	------	---------------

No objects
You don't have any objects in this bucket.

 Upload

Set Up Kinesis Data Stream and Streaming through Lambda

Create a Kinesis Data Stream

- Go to the Amazon Kinesis Console
- Go to Data Streams.
- Create a new stream name baseball-stream
- Leave other setting as default
- Your Kinesis Data Stream should look as following:

Amazon Kinesis > Data streams > baseball-stream

baseball-stream [Info](#)

[Delete](#)

Data stream summary

Status Active	Capacity mode On-demand Data retention period 1 day	ARN arn:aws:kinesis:us-west-1:124355677069:stream/baseball-stream	Creation time December 07, 2024 at 10:45 PST
------------------	--	--	---

Applications | Monitoring | **Configuration** | Enhanced fan-out (0) | Data viewer | Data analytics - new | Data stream sharing | EventBridge Pipes

Data stream capacity [Info](#)

[Edit capacity mode](#)

Capacity mode
On-demand

Tags - optional [Info](#)

[Manage tags](#)

Key	Value
No tags	
No tags associated with this stream	

[Manage tags](#)

Create a Lambda Function

- Navigate to the Lambda Console
- Create Function.
- Select "Author from scratch".
 - Function Name: stream_baseball_to_kinesis.
 - Runtime: Python 3.13.
 - Execution Role: Use the IAM role data engineer
 - Create function
 - Write the Lambda Function Code:
Below is the Python example for the Lambda function:
 - Click Deploy button after you finish inputting code

```
import boto3
import csv
import json
import os

# Initialize the Kinesis client
kinesis_client = boto3.client('kinesis')

def lambda_handler(event, context):
    try:
        # Get S3 bucket and object key from the event
        bucket_name = event['Records'][0]['s3']['bucket']['name']
        object_key = event['Records'][0]['s3']['object']['key']
```



```

print(f"Bucket: {bucket_name}, Key: {object_key}")

# Download the file from S3
s3_client = boto3.client('s3')
response = s3_client.get_object(Bucket=bucket_name,
Key=object_key)
data = response['Body'].read().decode('utf-8-sig').splitlines()

# Parse the CSV file
reader = csv.DictReader(data)

# Stream records to Kinesis
for row in reader:
    #skip row for no player and print it out for audit
    if not row['Player']:
        print(f"Skipping empty row: {row}")
        continue

    # Convert each row to JSON
    record = json.dumps(row)
    print(f"Sending record: {record}")

    # Send record to Kinesis Data Stream
    partition_key = row['Player']
    kinesis_client.put_record(
        StreamName='baseball-stream',
        Data=record,
        PartitionKey=partition_key
    )

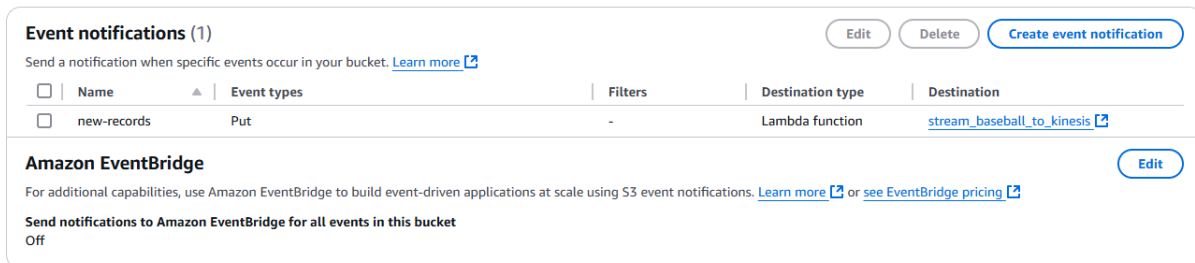
#verify stream complete
print("Streaming to Kinesis completed.")

#stream fail and reason for failing
except Exception as e:
    print(f"Streaming to Kinesis failed: {str(e)}")

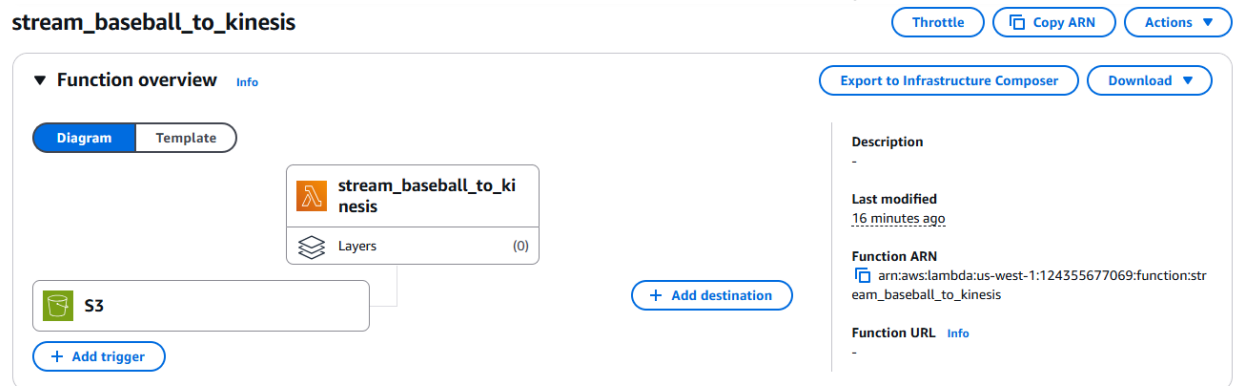
```

Configure S3 Event Notifications

- Navigate to the S3 bucket in the AWS Management Console.
- Go to the Properties tab and Event Notifications.
- Create a notification as following:
 - Event name: new-records
 - Event Type: PUT (for new file uploads).
 - Destination: Lambda Function (created in the next step).
 - Lambda function: stream_baseball_to_kinesis.
- Your Kinesis S3 Event Notifications should look as following:



- Your Lambda connected to S3 Notification should look as following:



Create a Lambda Test Function

The csv file of the game log statistics for lambda testing can be found in my github repository link:
https://github.com/Protprommart/LA-Dodgers-Cloud-ETL/blob/main/baseball_players_lambda_test.csv
Upload this csv file into baseball-statistics-dodgers S3 bucket as following:

baseball-statistics-dodgers [Info](#)

[Objects](#) | [Properties](#) | [Permissions](#) | [Metrics](#) | [Management](#) | [Access Points](#)

Objects (2) [Info](#)

[Copy S3 URI](#)[Copy URL](#)[Download](#)[Open](#)[Delete](#)[Actions](#)[Create folder](#)[Upload](#)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

< 1 >

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	baseball_players_march_to_May.csv	csv	December 7, 2024, 11:31:53 (UTC-08:00)	13.4 KB	Standard
<input type="checkbox"/>	baseball_players_lambda_test.csv	csv	December 7, 2024, 12:25:51 (UTC-08:00)	458.0 B	Standard

baseball_players_lambda_test.csv [Info](#)

[Copy S3 URI](#)[Download](#)[Open](#)[Object actions](#)

[Properties](#) | [Permissions](#) | [Versions](#)

Object overview

Owner

wareprotprom

AWS Region

US West (N. California) us-west-1

Last modified

December 7, 2024, 12:25:51 (UTC-08:00)

Size

458.0 B

Type

csv

Key

[baseball_players_lambda_test.csv](#)

S3 URI

s3://baseball-statistics-dodgers/baseball_players_lambda_test.csv

Amazon Resource Name (ARN)

arn:aws:s3::baseball-statistics-dodgers/baseball_players_lambda_test.csv

Entity tag (Etag)

[fd8f971d2fda89f275f78635e5fc12a9](#)

Object URL

https://baseball-statistics-dodgers.s3.us-west-1.amazonaws.com/baseball_players_lambda_test.csv

- Click on test button in stream_baseball_to_kinesis lambda function
- Name the test as teststream
- Use the S3 PUT event template.
- Update the event JSON to include:
 - Correct bucket name: baseball-statistics-dodgers.
 - File name (key): baseball_players_lambda_test.csv
 - Put in your awsRegion you are operating from
 - Put in the correct Etag of the object by getting it from S3 bucket object.
- Your test event JSON should look as following:

```
{
  "Records": [
    {
      "eventVersion": "2.1",
      "eventSource": "aws:s3",
      "awsRegion": "us-west-1",
      "eventTime": "2023-03-10T15:21:30.000Z",
```

```

    "eventName": "ObjectCreated:Put",
    "s3": {
      "s3SchemaVersion": "1.0",
      "bucket": {
        "name": "baseball-statistics-dodgers",
        "arn": "arn:aws:s3:::baseball-statistics-dodgers"
      },
      "object": {
        "key": "baseball_players_lambda_test.csv",
        "size": 801.0,
        "eTag": "6ced031204f54ad963f10e25556936d0",
        "sequencer": "00123456789abcdef"
      }
    }
  }
]
}

```

- Run the test code. Your output should be looking as following:

```

Status: Succeeded
Test Event Name: teststream

Response:
null

Function Logs:
START RequestId: 9f88412b-celf-487b-bf63-863aca0c49ee Version: $LATEST
Bucket: baseball-statistics-dodgers, Key:
baseball_players_lambda_test.csv
Sending record: {"Player": "Shohei Ohtani", "Team": "20-Mar", "OPP": "@
SD", "At Bats": "5", "Runs": "0", "Hits": "2", "Total Bases": "2", "2B":
"0", "3B": "0", "Home Runs": "0", "RunsBattedIn": "1", "Bases on Balls":
"0", "Intentional Walks": "0", "StrikeOut": "0", "StolenBases": "1",
"CaughtStealing": "0", "AVG": "0.4", "On-Base Percentage": "0.4",
"Slugging Percentage": "0.4", "HitByPitch": "0", "sacrifice bunt": "0",
"Sacrifice Flies": "0"}
Sending record: {"Player": "Shohei Ohtani", "Team": "21-Mar", "OPP": "vs
SD", "At Bats": "5", "Runs": "1", "Hits": "1", "Total Bases": "1", "2B":

```

```
"0", "3B": "0", "Home Runs": "0", "RunsBattedIn": "1", "Bases on Balls":
"0", "Intentional Walks": "0", "StrikeOut": "0", "StolenBases": "0",
"CaughtStealing": "0", "AVG": "0.3", "On-Base Percentage": "0.273",
"Slugging Percentage": "0.3", "HitByPitch": "0", "sacrifice bunt": "0",
"Sacrifice Flies": "1"}
```

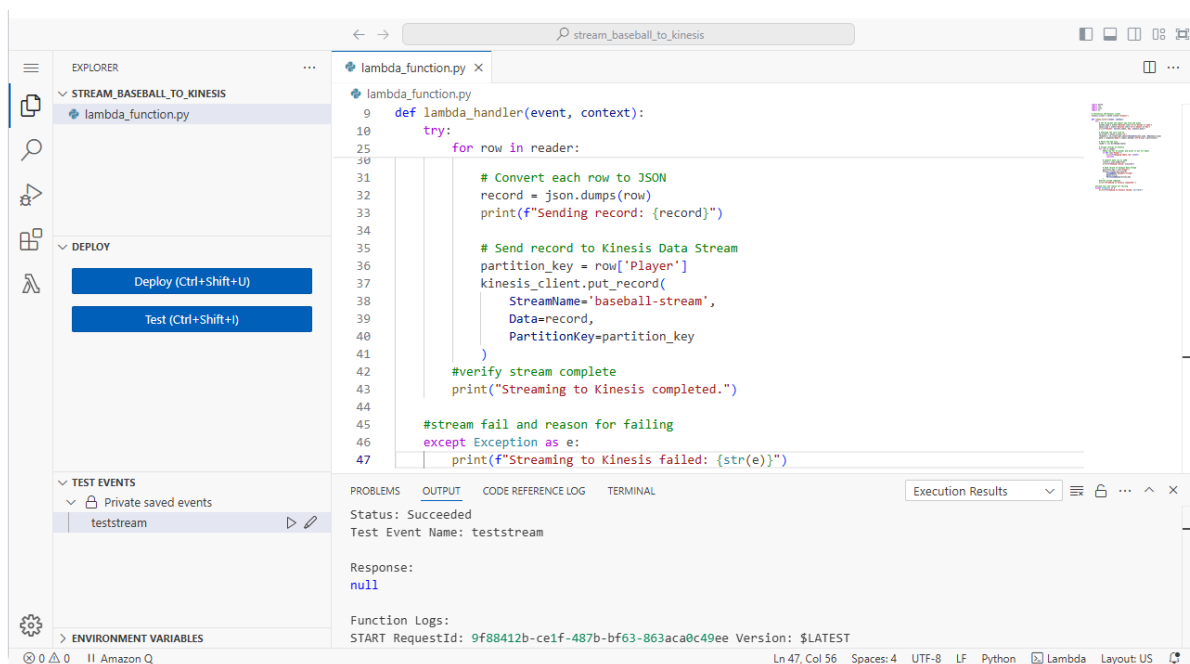
```
Sending record: {"Player": "Shohei Ohtani", "Team": "28-Mar", "OPP": "vs
STL", "At Bats": "3", "Runs": "1", "Hits": "2", "Total Bases": "3",
"2B": "1", "3B": "0", "Home Runs": "0", "RunsBattedIn": "0", "Bases on
Balls": "1", "Intentional Walks": "0", "StrikeOut": "1", "StolenBases":
"0", "CaughtStealing": "0", "AVG": "0.385", "On-Base Percentage": "0.4",
"Slugging Percentage": "0.462", "HitByPitch": "0", "sacrifice bunt":
"0", "Sacrifice Flies": "0"}
```

Streaming to Kinesis completed.

END RequestId: 9f88412b-ce1f-487b-bf63-863aca0c49ee

REPORT RequestId: 9f88412b-ce1f-487b-bf63-863aca0c49ee Duration:
2125.70 ms Billed Duration: 2126 ms Memory Size: 128 MB Max Memory
Used: 85 MB Init Duration: 538.16 ms

Request ID: 9f88412b-ce1f-487b-bf63-863aca0c49ee



Testing Streaming Integration Process

- Reupload baseball_players_march_to_may.csv data in baseball-statistics-dodgers S3 bucket

- Go to CloudWatch and navigate to Log groups under Logs tab
- Under Log groups find /aws/lambda/stream_baseball_to_kinesis to look at lambda streaming records
- In the log streams, click on the latest log stream

/aws/lambda/stream_baseball_to_kinesis

Actions View in Logs Insights Start tailing Search log group

▼ Log group details

Log class
Standard

ARN
arn:aws:logs:us-west-1:124355677069:log-group:/aws/lambda/stream_baseball_to_kinesis*

Creation time
8 hours ago

Retention
Never expire

Stored bytes
-

Metric filters
0

Subscription filters
0

Contributor Insights rules
-

KMS key ID
-

Anomaly detection
Configure

Data protection
-

Sensitive data count
-

Field indexes
Configure

Transformer
Configure

Log streams Tags Anomaly detection Metric filters Subscription filters Contributor Insights Data protection Field indexes - new Transformer - new

Log streams (17)

Filter log streams or try prefix search

Exact match Show expired

1

▼

2024/12/08/[\$LATEST]d8e5657966c742fa9c9ee0cc7f3e9d80

2024-12-08 03:16:53 (UTC)

- Check if the stream went successfully. Your log streams should look as following:

Log events

You can use the filter bar below to search for and match terms, phrases, or values in your log events.

Filter events - press enter to search

Clear 1m 30m 1h 12h Custom UTC timezone Display

Timestamp	Message
There are older events to load.	
2024-12-08T03:16:52.339Z	Sending record: {"Player": "Freddie Freeman", "Team": "8-Apr", "OPP": "0 MIN", "At Bats": "4", "Runs": "0", "Hits": "2", "Total Bases": "2", "2B": "0", "3B": "0", "Home Runs": "0", "RunsBat...
2024-12-08T03:16:52.359Z	Sending record: {"Player": "Freddie Freeman", "Team": "9-Apr", "OPP": "0 MIN", "At Bats": "5", "Runs": "1", "Hits": "1", "Total Bases": "1", "2B": "0", "3B": "0", "Home Runs": "0", "RunsBat...
2024-12-08T03:16:52.379Z	Sending record: {"Player": "Freddie Freeman", "Team": "10-Apr", "OPP": "0 MIN", "At Bats": "3", "Runs": "0", "Hits": "1", "Total Bases": "2", "2B": "1", "3B": "0", "Home Runs": "0", "RunsBa...
2024-12-08T03:16:52.399Z	Sending record: {"Player": "Freddie Freeman", "Team": "12-Apr", "OPP": "vs SD", "At Bats": "5", "Runs": "1", "Hits": "1", "Total Bases": "3", "2B": "0", "3B": "0", "Home Runs": "0", "RunsBa...
2024-12-08T03:16:52.419Z	Sending record: {"Player": "Freddie Freeman", "Team": "13-Apr", "OPP": "vs SD", "At Bats": "4", "Runs": "0", "Hits": "0", "Total Bases": "0", "2B": "0", "3B": "0", "Home Runs": "0", "RunsBa...
2024-12-08T03:16:52.439Z	Sending record: {"Player": "Freddie Freeman", "Team": "14-Apr", "OPP": "vs SD", "At Bats": "4", "Runs": "1", "Hits": "1", "Total Bases": "2", "2B": "1", "3B": "0", "Home Runs": "0", "RunsBa...
2024-12-08T03:17:59.949Z	Sending record: {"Player": "Shohei Ohtani", "Team": "20-Mar", "OPP": "0 SD", "At Bats": "5", "Runs": "0", "Hits": "2", "Total Bases": "2", "2B": "0", "3B": "0", "Home Runs": "0", "RunsBatte...
2024-12-08T03:18:00.301Z	Sending record: {"Player": "Shohei Ohtani", "Team": "21-Mar", "OPP": "vs SD", "At Bats": "5", "Runs": "1", "Hits": "1", "Total Bases": "1", "2B": "0", "3B": "0", "Home Runs": "0", "RunsBett...
2024-12-08T03:18:00.320Z	Sending record: {"Player": "Shohei Ohtani", "Team": "28-Mar", "OPP": "vs STL", "At Bats": "3", "Runs": "1", "Hits": "2", "Total Bases": "3", "2B": "1", "3B": "0", "Home Runs": "0", "RunsBat...
2024-12-08T03:18:00.340Z	Sending record: {"Player": "Shohei Ohtani", "Team": "29-Mar", "OPP": "vs STL", "At Bats": "4", "Runs": "0", "Hits": "0", "Total Bases": "0", "2B": "0", "3B": "0", "Home Runs": "0", "RunsBat...
2024-12-08T03:18:00.380Z	Sending record: {"Player": "Shohei Ohtani", "Team": "30-Mar", "OPP": "vs STL", "At Bats": "5", "Runs": "1", "Hits": "1", "Total Bases": "1", "2B": "0", "3B": "0", "Home Runs": "0", "RunsBat...
2024-12-08T03:18:00.400Z	Sending record: {"Player": "Shohei Ohtani", "Team": "31-Mar", "OPP": "vs STL", "At Bats": "4", "Runs": "1", "Hits": "1", "Total Bases": "2", "2B": "1", "3B": "0", "Home Runs": "0", "RunsBat...
2024-12-08T03:18:00.420Z	Sending record: {"Player": "Mookie Betts", "Team": "28-Mar", "OPP": "0 SD", "At Bats": "4", "Runs": "0", "Hits": "2", "Total Bases": "2", "2B": "0", "3B": "0", "Home Runs": "0", "RunsBatted...
2024-12-08T03:18:01.280Z	Sending record: {"Player": "Mookie Betts", "Team": "1-Apr", "OPP": "vs SF", "At Bats": "3", "Runs": "3", "Hits": "2", "Total Bases": "5", "2B": "1", "3B": "1", "Home Runs": "0", "RunsBatted...
2024-12-08T03:18:01.390Z	Sending record: {"Player": "Mookie Betts", "Team": "2-Apr", "OPP": "vs SF", "At Bats": "5", "Runs": "2", "Hits": "2", "Total Bases": "5", "2B": "0", "3B": "0", "Home Runs": "1", "RunsBatted...
2024-12-08T03:18:01.320Z	Sending record: {"Player": "Mookie Betts", "Team": "3-Apr", "OPP": "vs SF", "At Bats": "3", "Runs": "0", "Hits": "1", "Total Bases": "1", "2B": "0", "3B": "0", "Home Runs": "0", "RunsBatted...
2024-12-08T03:18:01.341Z	Sending record: {"Player": "Mookie Betts", "Team": "5-Apr", "OPP": "0 CHC", "At Bats": "5", "Runs": "1", "Hits": "1", "Total Bases": "1", "2B": "0", "3B": "0", "Home Runs": "0", "RunsBatted...
2024-12-08T03:18:01.360Z	Sending record: {"Player": "Mookie Betts", "Team": "6-Apr", "OPP": "0 CHC", "At Bats": "3", "Runs": "0", "Hits": "0", "Total Bases": "0", "2B": "0", "3B": "0", "Home Runs": "0", "RunsBatted...
2024-12-08T03:18:01.400Z	Sending record: {"Player": "Mookie Betts", "Team": "7-Apr", "OPP": "0 CHC", "At Bats": "3", "Runs": "1", "Hits": "0", "Total Bases": "0", "2B": "0", "3B": "0", "Home Runs": "0", "RunsBatted...

You should be able to see how lambda is streaming to kinesis by organizing according to the baseball player name. This is because we specify in the code that we are going to use the player column as a partition key.

- Go to baseball-stream Kinesis data stream and you can check records coming in to Kinesis under the data viewer tab.
- View Shardid-0000000000001 for Freddie Freeman, 0000000000002 for Shohei Ohtani, 0000000000003 for Mookie Betts

Shardid0000000000001

Shard		Starting position Info	
shardid-000000000001		Trim horizon	
		Get records	

Records (24)			Next records
Shard: shardid-000000000001 Starting position: Trim horizon			
<input type="text" value="Find records"/>			
<div> <div>< 1 ></div> <div>⚙</div> </div>			
Partition key	Data	Approximate arrival timestamp	Sequence number
Freddie Freeman	{"Player": "Freddie Freeman", "Team": "20-Mar", "OPP": "@ SD", "At Bats": "2", ...	December 07, 2024 at 12:11:07 PST	49658425433342475785413592021030895506600116280799264786
Freddie Freeman	{"Player": "Freddie Freeman", "Team": "21-Mar", "OPP": "vs SD", "At Bats": "4", ...	December 07, 2024 at 12:11:07 PST	49658425433342475785413592021032104432419730909973970962
Freddie Freeman	{"Player": "Freddie Freeman", "Team": "28-Mar", "OPP": "vs STL", "At Bats": "3", ...	December 07, 2024 at 12:11:07 PST	4965842543334247578541359202103313358239345539148677138
Freddie Freeman	{"Player": "Freddie Freeman", "Team": "29-Mar", "OPP": "vs STL", "At Bats": "4", ...	December 07, 2024 at 12:11:07 PST	49658425433342475785413592021034522284058960168323383314
Freddie Freeman	{"Player": "Freddie Freeman", "Team": "30-Mar", "OPP": "vs STL", "At Bats": "5", ...	December 07, 2024 at 12:11:07 PST	49658425433342475785413592021035731209878574797498089490

Shardid000000000002

Applications	Monitoring	Configuration	Enhanced fan-out (0)	Data viewer	Data analytics - <i>new</i>	Data stream sharing	EventBridge Pipes
--------------	------------	---------------	----------------------	--------------------	-----------------------------	---------------------	-------------------

Shard		Starting position Info	
shardid-0000000000002		Trim horizon	
		Get records	

Records (27)			Next records
Shard: shardid-0000000000002 Starting position: Trim horizon			
<input type="text" value="Find records"/>			
<div> <div>< 1 ></div> <div>⚙</div> </div>			
Partition key	Data	Approximate arrival timestamp	Sequence number
Shohei Ohtani	{"Player": "Shohei Ohtani", "Team": "20-Mar", "OPP": "@ SD", "At Bats": "5", ...	December 07, 2024 at 12:11:07 PST	49658425433364776530612122644157924115037389023489294370
Shohei Ohtani	{"Player": "Shohei Ohtani", "Team": "21-Mar", "OPP": "vs SD", "At Bats": "5", ...	December 07, 2024 at 12:11:07 PST	49658425433364776530612122644159133040857003652664000546
Shohei Ohtani	{"Player": "Shohei Ohtani", "Team": "28-Mar", "OPP": "vs STL", "At Bats": "3", ...	December 07, 2024 at 12:11:07 PST	49658425433364776530612122644160341966676618281838706722
Shohei Ohtani	{"Player": "Shohei Ohtani", "Team": "29-Mar", "OPP": "vs STL", "At Bats": "4", ...	December 07, 2024 at 12:11:07 PST	49658425433364776530612122644161550892496232979732889634

Shardid000000000003

Applications	Monitoring	Configuration	Enhanced fan-out (0)	Data viewer	Data analytics - <i>new</i>	Data stream sharing	EventBridge Pipes
--------------	------------	---------------	----------------------	--------------------	-----------------------------	---------------------	-------------------

Shard		Starting position Info	
shardid-0000000000003		Trim horizon	
		Get records	

Records (21)			Next records
Shard: shardid-0000000000003 Starting position: Trim horizon			
<input type="text" value="Find records"/>			
<div> <div>< 1 ></div> <div>⚙</div> </div>			
Partition key	Data	Approximate arrival timestamp	Sequence number
Mookie Betts	{"Player": "Mookie Betts", "Team": "20-Mar", "OPP": "@ SD", "At Bats": "4", "R...",	December 07, 2024 at 12:11:07 PST	49658425433387077275810653267306713388227725228762988594
Mookie Betts	{"Player": "Mookie Betts", "Team": "21-Mar", "OPP": "vs SD", "At Bats": "5", "R...",	December 07, 2024 at 12:11:07 PST	49658425433387077275810653267307922314047339857937694770
Mookie Betts	{"Player": "Mookie Betts", "Team": "28-Mar", "OPP": "vs STL", "At Bats": "2", "R...",	December 07, 2024 at 12:11:07 PST	49658425433387077275810653267309131239866954487112400946
Mookie Betts	{"Player": "Mookie Betts", "Team": "29-Mar", "OPP": "vs STL", "At Bats": "3", "R...",	December 07, 2024 at 12:11:07 PST	49658425433387077275810653267310340165686569116287107122

As you can see here that each player's statistics are being separated into each shards.

Create a Data Warehouse in Redshift

- Go to Redshift Serverless
- Name target namespace as baseball-data-warehouse. The namespace will host the database
- Name workgroup as workgroupladodgers. Workgroup will host the server
- In VPC choose ITL-vpc that we created
- Choose security group that create as a default with the vpc we created
- Choose private subnets we created for subnet
- Note: there needs to be at least 37 free IP addresses in 2 subnets. Each subnet should be in a different Availability Zone.
- Add data-engineer IAM roles to it.

You should have the following Namespace and workgroup.

baseball-data-warehouse [Info](#)

[Refresh](#) [Actions](#) [Query data](#)

General information

Namespace baseball-data-warehouse	Status Available	Admin user name admin
Namespace ID 3b3d1d99-c7b5-4316-8cf1-28ff4edf8f9a	Date created December 03, 2024, 22:02 (UTC-08:00)	Database name dev
Namespace ARN arn:aws:redshift-serverless:us-west-1:124355677069:namespace/3b3d1d99-c7b5-4316-8cf1-28ff4edf8f9a	Storage used 811 MB	Total table count -

Workgroup | Data backup | Database | **Security and encryption** | Datashares | Zero-ETL integrations | Resource policy | Tags

Permissions [Info](#)

[Manage IAM roles](#)

IAM roles	Status	Amazon Resource Name (ARN)
data-engineer	in-sync	arn:aws:iam::124355677069:role/data-engineer



This workgroup is associated with a namespace

To manage database objects and users, navigate to the namespace that this workgroup is associated with.

baseball-data-warehouse [X](#)

workgroupladodgers [Info](#)

[Actions](#) [Query data](#)

General information

[Refresh](#)

Workgroup workgroupladodgers	Date created December 03, 2024, 22:02 (UTC-08:00)	Endpoint workgroupladodgers.124355677069.us-west-1.redshift-serverless.amazonaws.com:5439/dev
Namespace baseball-data-warehouse	Status Available	JDBC URL jdbc:redshift://workgroupladodgers.124355677069.us-west-1.redshift-serverless.amazonaws.com:5439/dev
Workgroup ARN arn:aws:redshift-serverless:us-west-1:124355677069:workgroup/f6ea092c-f0df-4456-8224-5451bdbb471d	Base capacity 128 RPU	ODBC URL Driver={Amazon Redshift (x64)}; Server=workgroupladodgers.124355677069.us-west-1.redshift-serverless.amazonaws.com; Database=dev
Workgroup version 1.0.79237	Custom domain name -	
	Patch version Patch 186	

Transforming Data with AWS Glue

Create Database

- Go to Amazon Glue and click on Databases
- Create new database by click add database

- Enter the name of database baseball-stats
- Click create database

Your Database should look as the following:

baseball-stats

Last updated (UTC)
 December 4, 2024 at 23:52:29

Edit

Delete

Database properties

Name baseball-stats	Description -	Location -	Created on (UTC) December 4, 2024 at 23:52:27
-------------------------------	-------------------------	----------------------	---

Tables (0)

Last updated (UTC)
 December 4, 2024 at 23:52:29

Delete

Add tables using crawler

Add table

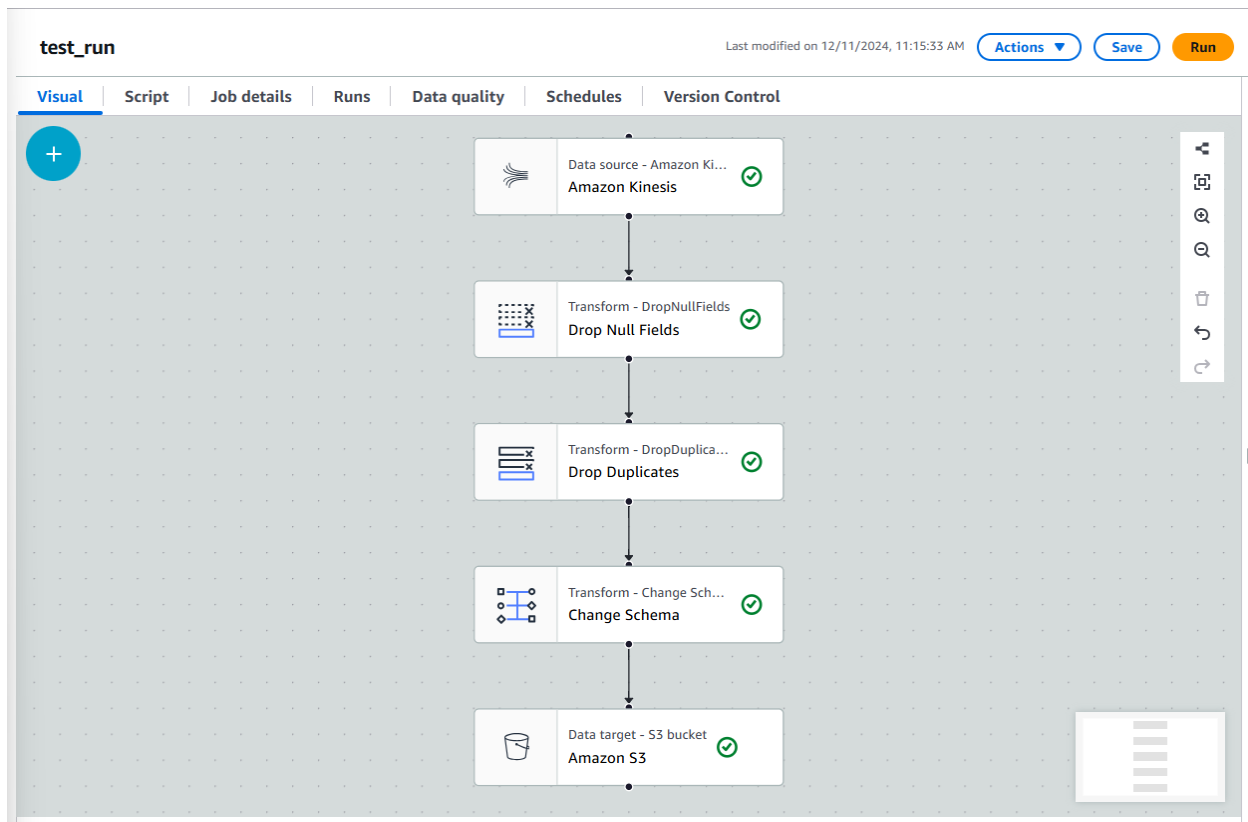
View and manage all available tables.

< 1 >

Name	Database	Location	Classific...	Depreca...	View data	Data quality	Column statistics
No available tables							

Use Visual ETL to transform data

- Go to AWS Glue and to Visual ETL
- Choose the source as baseball-stream Kinesis stream where we have untransformed data
 - Choose stream details
 - Starting position: Earliest
 - Window size:100
 - Choose dataformat JSON
- Choose Drop Null Fields
 - For dropnullfields, click check for empty string, "null" string, and -1 Integer
- Choose Drop Duplicates
- Choose change schema for transformation
 - keep only player, team change to date, homeruns, runbattedin, avg, and on-base percentage
 - Cast homeruns and runbattedin to integer while cast avg and on-base percentage for double
- Choose target as S3 bucket clean-baseball-statistics to store transformed data
 - Choose JSON format, uncompressed
 - Uncheck the box for data quality because it will failed the run job for not importing aws data quality library.
- Go to job detail tab
 - Make sure the IAM role is data engineer
 - Make sure type is spark streaming
 - Choose worker type G.25X which is suitable for streaming
 - Check automatically scale the number of workers
 - Input maximum number of workers you want to use. I chose 5 workers



Refs:

<https://www.youtube.com/watch?v=6ggTFOfUxU>










<https://www.youtube.com/watch?v=ziEEeL1egPk>

- Check in clean-baseball-statistics S3 bucket
- You should see folder inside called ingest_year=2024/, ingest_month=12/, ingest_day=8/, and hour of ingestion

ingest_hour=05/

Copy S3 URI

Objects Properties

Objects (8) Info   Copy S3 URI  Copy URL  Download  Open  Delete  Actions  Create folder  Upload

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	run-1733635435731-part-r-00000	-	December 7, 2024, 21:23:59 (UTC-08:00)	0 B	Standard
<input type="checkbox"/>	run-1733635435731-part-r-00001	-	December 7, 2024, 21:23:59 (UTC-08:00)	2.7 KB	Standard
<input type="checkbox"/>	run-1733635435731-part-r-00002	-	December 7, 2024, 21:23:59 (UTC-08:00)	2.9 KB	Standard
<input type="checkbox"/>	run-1733635435731-part-r-00003	-	December 7, 2024, 21:23:59 (UTC-08:00)	2.1 KB	Standard
<input type="checkbox"/>	run-1733635442529-part-r-00000	-	December 7, 2024, 21:24:03 (UTC-08:00)	0 B	Standard
<input type="checkbox"/>	run-1733635442529-part-r-00001	-	December 7, 2024, 21:24:04 (UTC-08:00)	20.7 KB	Standard
<input type="checkbox"/>	run-1733635442529-part-r-00002	-	December 7, 2024, 21:24:03 (UTC-08:00)	5.5 KB	Standard
<input type="checkbox"/>	run-1733635442529-part-r-00003	-	December 7, 2024, 21:24:03 (UTC-08:00)	12.3 KB	Standard

Use crawler to extract the schema and metadata of transformed data

- Go to AWS Glue and to crawler
- Click create crawler
- Name the crawler baseball-crawler
- Add datasource and choose S3 bucket clean-baseball-statistics
- choose data engineer IAM role
- Choose baseball database
- Choose prefix ladodger_stream_
- Choose on demand for crawler schedule

baseball-crawler

Last updated (UTC) December 8, 2024 at 18:50:32  Run crawler  Edit  Delete

Crawler properties

Name baseball-crawler	IAM role data-engineer	Database baseball-stats	State READY
Description -	Security configuration -	Lake Formation configuration -	Table prefix ladodger_stream_
Maximum table threshold -			

Advanced settings

Crawler runs Schedule Data sources Classifiers Tags

Crawler runs (0)

The list of crawler runs for this crawler.

Filter data

Filter by a date and time range

Start time (UTC)	End time (UTC)	Current/last duration	Status	DPU hours	Table changes
You don't have any crawler runs.					

Run crawler

After that you should be able to see a table in the baseball-stream database once the crawler is done.

Crawler runs | Schedule | Data sources | Classifiers | Tags

Crawler runs (1)
The list of crawler runs for this crawler.

Filter data

Filter by a date and time range

Start time (UTC) | End time (UTC) | Current/last duration | Status | DPU hours | Table changes

December 8, 2024 at 18:51:11 | December 8, 2024 at 18:52:05 | 53 s | Completed | 0.083 | 1 table change, 1 partition change

Stop run | View CloudWatch logs | View run details

Go to database baseball-stats and to ladodger_stream_clean_baseball_statistics that we made

ladodger_stream_clean_baseball_statistics

Last updated (UTC) | Version 1 (Current version) | Actions

Table overview | Data quality - new

Table details

Name ladodger_stream_clean_baseball_statistics Database baseball-stats Description - Last updated December 8, 2024 at 18:52:04	Classification JSON Location s3://clean-baseball-statistics/ Connection -	Deprecated - Column statistics No statistics
---	---	---

Advanced properties

Schema | Partitions | Indexes | Column statistics - new

Schema (10)
View and manage the table schema.

Filter schemas

#	Column name	Data type	Partition key	Comment
1	player	string	-	-
2	date	string	-	-
3	home runs	int	-	-
4	runsbatte din	int	-	-
5	avg	double	-	-
6	on-base percentage	double	-	-
7	ingest_year	string	Partition (0)	-
8	ingest_month	string	Partition (1)	-
9	ingest_day	string	Partition (2)	-
10	ingest_hour	string	Partition (3)	-

Edit schema as JSON | Edit schema

We see that ingest year, month, day, and hour column were added to save when records were brought in by Kinesis

Use Redshift Query Editor to visualize and query transformed data

Query Transformed Data

- Go to workgroup on redshift
- Click on query data
- In redshift query editor go to external databases, and into baseball-stats database, and into the ladodger_stream_clean_baseball_statistics table created by crawler
- Right click on the table, choose select table

The screenshot shows the AWS Redshift Query Editor v2 interface. The left sidebar displays the 'Queries' tab with a tree view of databases and tables. The 'baseball-stats' database is expanded, showing the 'ladodger_stream_clean_baseball_statistics' table. The main editor area shows a SQL query: `SELECT * FROM \"awsdatacatalog\".\"baseball-stats\".\"ladodger_stream_clean_baseball_statistics\";`. Below the editor, a table preview is displayed with columns: player, date, home runs, runsbattedin, avg, and on-base percentage. The table contains 100 rows of data for 'Mookie Betts'.

Your table should display as this.

Visualize Data

Our goal in querying and visualizing is, let's say that you have to deliver best players statistics to the company that carries out a sports betting service, for the company to give statistics to users whether they are going to bet on among three batters. We are going to be looking especially at 2 statistics for batters position including an average of, on-base percentage, on average what is the percentage that the player hit the ball and make it safe to a base, and total home runs they made in this season so far.

- Use the following query in redshift query editor to calculate statistics

```
SELECT
    "player",
    SUM("home runs") AS "total_homeruns",
    AVG("on-base percentage") AS "average_on_base_percent"
FROM
    "awsdatacatalog"."baseball-
stats"."ladodger_stream_clean_baseball_statistics"
GROUP BY
    "player";
```

- You should be able to see the following query result

Run

Limit 100

Explain

Isolated session

Serverless: w...

dev

```


1 SELECT
2     "player",
3     SUM("home runs") AS "total_homeruns",
4     AVG("on-base percentage") AS "average_on_base_percent"
5 FROM
6     "awsdatacatalog"."baseball-stats"."ladodger_stream_clean_baseball_statistics"
7 GROUP BY
8     "player";

```

Result 1 (3)

	player	total_homeruns	average_on_base_per...
<input type="checkbox"/>	Mookie Betts	5	0.6578889
<input type="checkbox"/>	Freddie Freeman	2	0.45421875
<input type="checkbox"/>	Shohei Ohtani	0	0.3365

- Click on chart button to view query result in chart
- Choose the following:
 - type: barchart
 - Values: total homeruns
 - Labels: player
- Choose values as average_on_base_percent to view horizontal bar-chart of average on-base percentage
- You can optionally export this chart by click on export button



- Based on these statistics, a batter that you should bet on is Mookie Betts because he has the highest total home runs and average on base percentage.

